

HW10

Lun Li.14415

2025-04-20

```
library(coda)

## Warning: package 'coda' was built under R version 4.3.3

library(rjags)

## Warning: package 'rjags' was built under R version 4.3.3

## Linked to JAGS 4.3.2

## Loaded modules: basemod,bugs

library(tidyr)
library(dplyr)

## 
## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':
## 
##     filter, lag

## The following objects are masked from 'package:base':
## 
##     intersect, setdiff, setequal, union

library(coda)
```

Problem 1

a

```
frd <- read.table('/Users/landlee/Desktop/STAT 3303/HW10/friends.txt',
                  header = TRUE)
```

```

# MCMC parameters
niter <- 30000
nburn <- 10000
nadapt <- 10000

# Number of schools and total observations
Ns <- length(unique(frd$school))
N <- nrow(frd)

# Prepare JAGS data
data_jags <- list(
  y = frd$y,
  x = frd$x,
  school = frd$school,
  Ns = Ns,
  N = N
)

# Initial values
set.seed(123)
inits <- list(
  mu_alpha = 0,
  mu_beta = 0,
  sigma_alpha = 1,
  sigma_beta = 1,
  alpha = rnorm(Ns, 0, 1),
  beta = rnorm(Ns, 0, 1)
)

model_string <- "model {
  for (i in 1:N) {
    y[i] ~ dbern(theta[i])
    logit(theta[i]) <- alpha[school[i]] + beta[school[i]] * x[i]
  }

  for (j in 1:Ns) {
    alpha[j] ~ dnorm(mu_alpha, tau_alpha)
    beta[j] ~ dnorm(mu_beta, tau_beta)
  }

  mu_alpha ~ dnorm(0, 1/9)
  mu_beta ~ dnorm(0, 1/9)
  sigma_alpha ~ dunif(0, 3)
  sigma_beta ~ dunif(0, 3)

  tau_alpha <- pow(sigma_alpha, -2)
  tau_beta <- pow(sigma_beta, -2)
}"
```

```

fit1 = jags.model(textConnection(model_string),
                   data = data_jags,
                   inits = inits,
                   n.chains = 1,
```

```

    n.adapt = nadapt)

## Compiling model graph
##   Resolving undeclared variables
##   Allocating nodes
## Graph information:
##   Observed stochastic nodes: 1520
##   Unobserved stochastic nodes: 20
##   Total graph size: 4639
##
## Initializing model

update(fit1, nburn) # Burn-in phase

fit.samples = coda.samples(fit1,
                           c("mu_alpha", "mu_beta", "sigma_alpha", "sigma_beta", "alpha", "beta"),
                           n.iter = niter)

summary(fit.samples)

##
## Iterations = 20001:50000
## Thinning interval = 1
## Number of chains = 1
## Sample size per chain = 30000
##
## 1. Empirical mean and standard deviation for each variable,
##    plus standard error of the mean:
##
##          Mean      SD  Naive SE Time-series SE
## alpha[1] -0.61638 0.2566 0.0014815      0.007778
## alpha[2] -0.20055 0.1619 0.0009345      0.001943
## alpha[3] -0.07552 0.1669 0.0009639      0.002359
## alpha[4] -0.07712 0.1649 0.0009520      0.002318
## alpha[5]  0.09777 0.1940 0.0011198      0.004531
## alpha[6] -0.11047 0.1652 0.0009535      0.001989
## alpha[7] -0.23545 0.1656 0.0009558      0.002279
## alpha[8] -0.21835 0.1649 0.0009520      0.002150
## beta[1]  -1.26382 0.4147 0.0023945      0.010226
## beta[2]  -0.38304 0.2352 0.0013578      0.002917
## beta[3]  -0.41145 0.2380 0.0013742      0.002886
## beta[4]  -0.40862 0.2358 0.0013612      0.003169
## beta[5]  -0.18133 0.2679 0.0015465      0.005922
## beta[6]  -0.77696 0.2500 0.0014434      0.003076
## beta[7]  -0.76385 0.2524 0.0014575      0.003574
## beta[8]  -0.33368 0.2390 0.0013799      0.003230
## mu_alpha -0.17986 0.1485 0.0008574      0.001679
## mu_beta  -0.56176 0.2302 0.0013292      0.002289
## sigma_alpha 0.31415 0.1819 0.0010500      0.005164
## sigma_beta  0.51230 0.2755 0.0015904      0.006636
##
## 2. Quantiles for each variable:

```

```

##          2.5%     25%     50%     75%   97.5%
## alpha[1] -1.11402 -0.79610 -0.61638 -0.4353529 -0.12936
## alpha[2] -0.52798 -0.30519 -0.19644 -0.0937009  0.10944
## alpha[3] -0.39485 -0.18791 -0.08178  0.0344720  0.26269
## alpha[4] -0.39011 -0.18803 -0.08304  0.0298063  0.25752
## alpha[5] -0.24503 -0.04358  0.09088  0.2324670  0.48675
## alpha[6] -0.42870 -0.21922 -0.11432 -0.0050599  0.23056
## alpha[7] -0.57545 -0.34332 -0.23008 -0.1250262  0.08054
## alpha[8] -0.55365 -0.32540 -0.21445 -0.1081875  0.09711
## beta[1] -2.12085 -1.54677 -1.24464 -0.9603454 -0.52941
## beta[2] -0.82718 -0.54264 -0.39173 -0.2267795  0.09682
## beta[3] -0.86836 -0.57374 -0.41844 -0.2516141  0.06539
## beta[4] -0.86532 -0.56757 -0.41396 -0.2520415  0.06630
## beta[5] -0.67261 -0.37233 -0.18618  0.0006705  0.35804
## beta[6] -1.28792 -0.94149 -0.76711 -0.6032845 -0.30751
## beta[7] -1.27974 -0.92862 -0.75429 -0.5913783 -0.29248
## beta[8] -0.78439 -0.49870 -0.33838 -0.1763076  0.14862
## mu_alpha -0.48478 -0.26359 -0.17848 -0.0936096  0.11113
## mu_beta  -1.03410 -0.69114 -0.55949 -0.4309090 -0.09974
## sigma_alpha 0.02787  0.19937  0.29020  0.3992051  0.72824
## sigma_beta  0.10276  0.33356  0.46639  0.6371775  1.19224

# Extract MCMC samples into matrix
mcmc_matrix <- as.matrix(fit.samples)

# Choose key parameters to visualize
params <- c("mu_alpha", "mu_beta", "sigma_alpha", "sigma_beta")

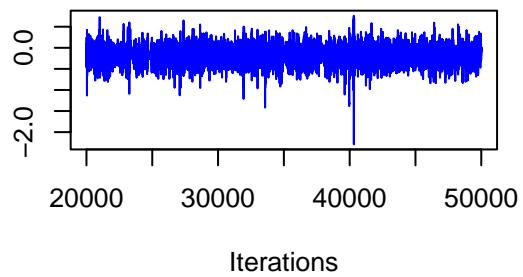
# Set up plotting layout: 2 rows, 2 columns (each parameter gets trace + density)
par(mfrow = c(2, 2)) # 2 plots per row

# Loop through parameters and plot
for (param in params) {
  traceplot(fit.samples[, param],
            main = paste("Traceplot of", param),
            col = "blue", lwd = 1)

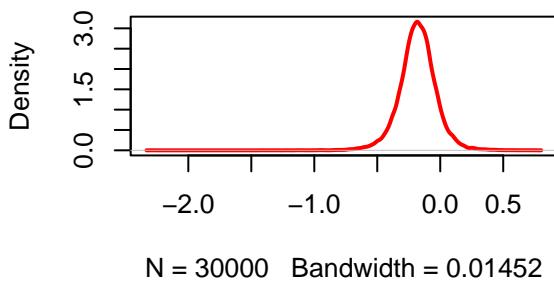
  plot(density(mcmc_matrix[, param]),
        main = paste("Posterior of", param),
        col = "red", lwd = 2)
}

```

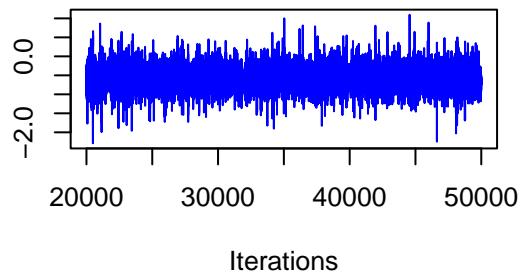
Traceplot of mu_alpha



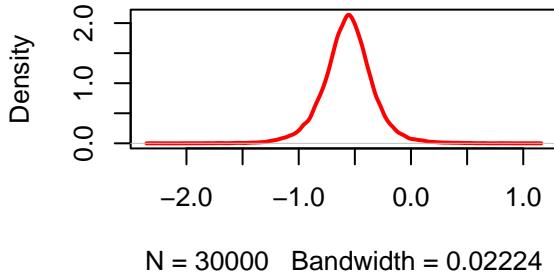
Posterior of mu_alpha

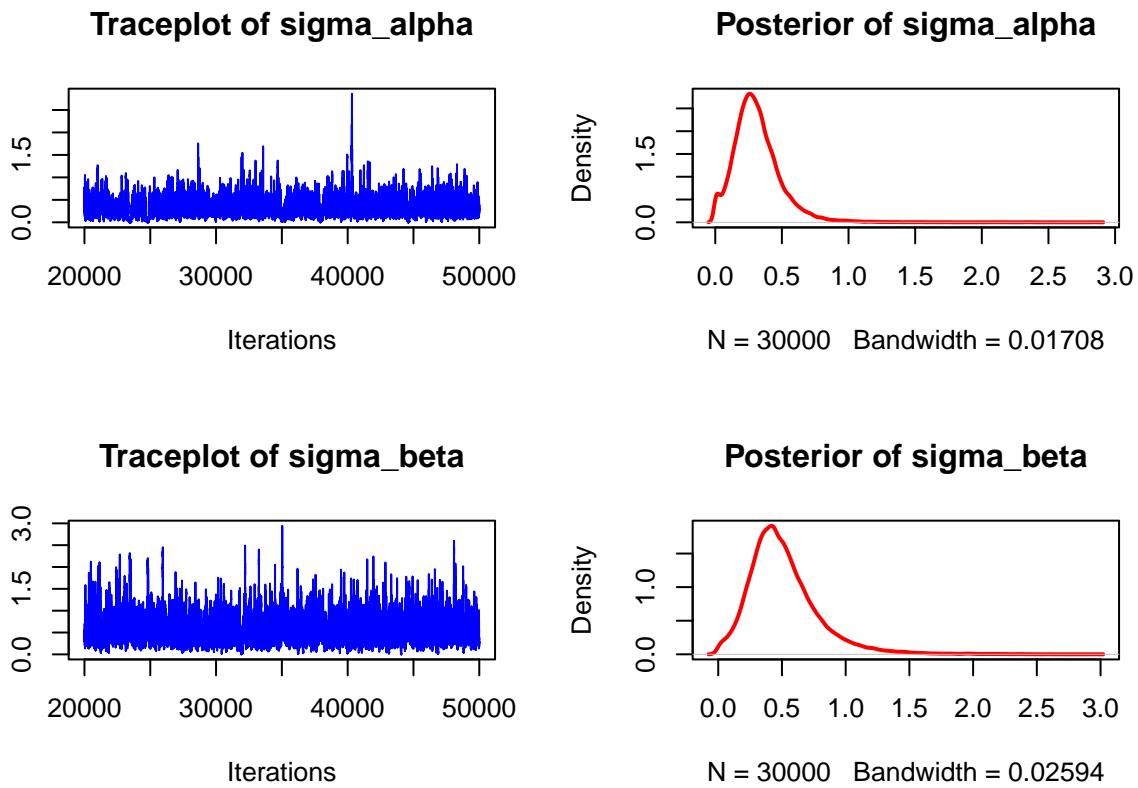


Traceplot of mu_beta



Posterior of mu_beta





```
fit.samples.adapted <- window(fit.samples, start = nadapt + 1, end = niter)
```

```
## Warning in FUN(X[[i]], ...): start value not changed
```

```
# ESS
ess <- effectiveSize(fit.samples.adapted)
cat("Effective Sample Size (ESS):\n")
```

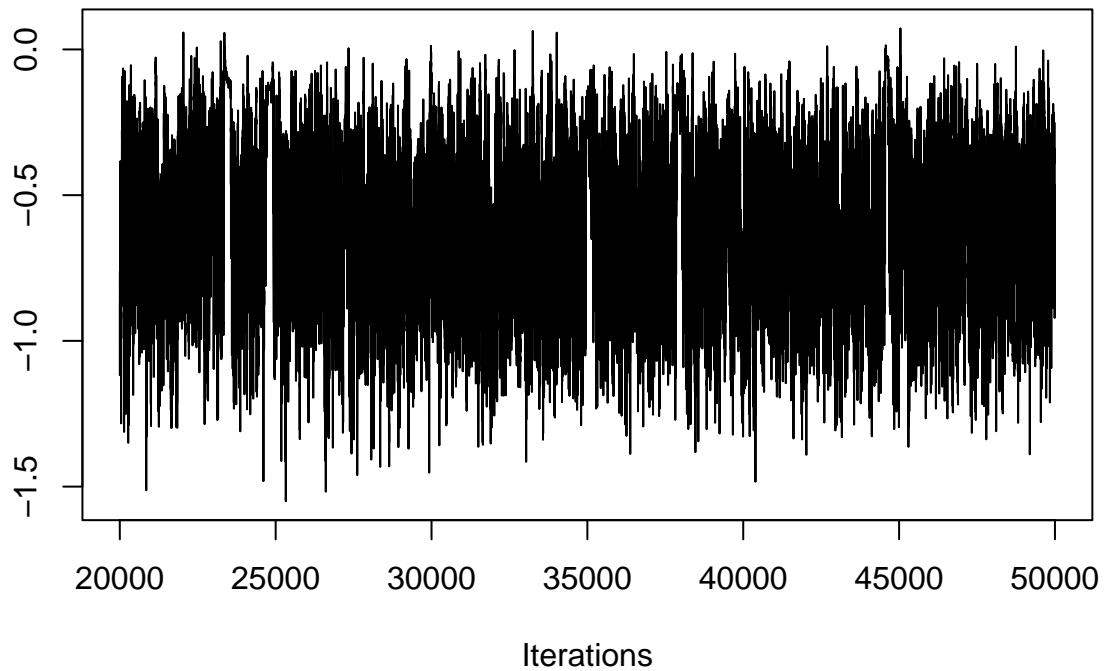
```
## Effective Sample Size (ESS):
```

```
print(round(ess, 2))
```

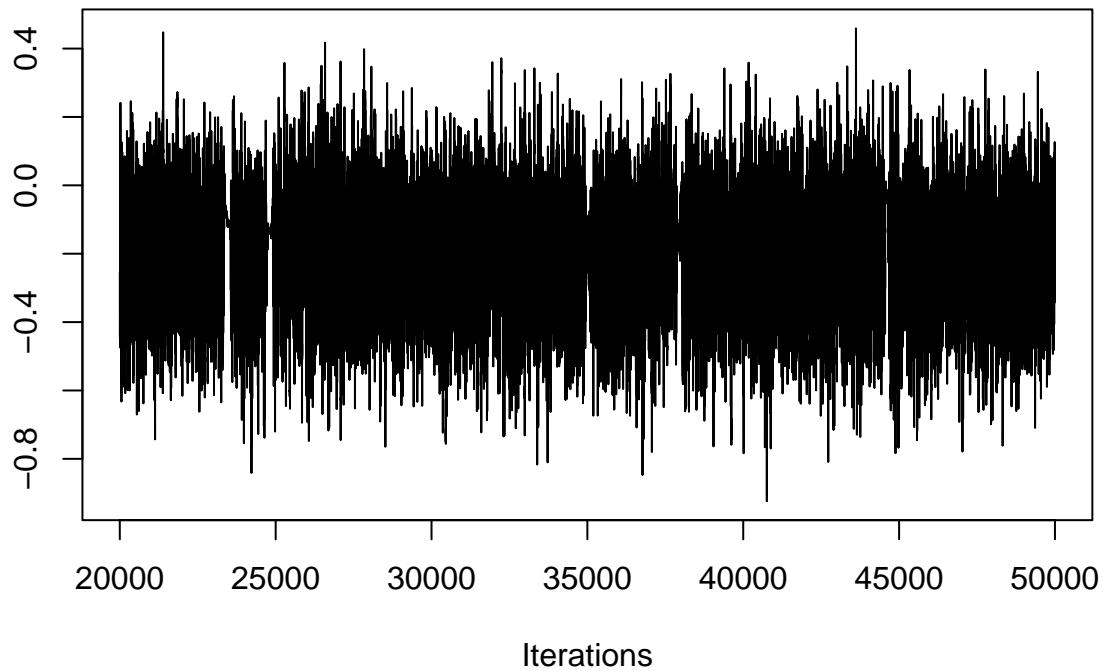
##	alpha[1]	alpha[2]	alpha[3]	alpha[4]	alpha[5]	alpha[6]
##	362.54	2222.50	1958.60	1881.29	769.51	2334.86
##	alpha[7]	alpha[8]	beta[1]	beta[2]	beta[3]	beta[4]
##	1600.65	1630.39	522.25	2237.81	2485.14	2046.31
##	beta[5]	beta[6]	beta[7]	beta[8]	mu_alpha	mu_beta
##	897.24	2379.70	1836.31	1632.58	2193.41	3458.82
##	sigma_alpha	sigma_beta				
##	493.17	609.02				

```
traceplot(fit.samples)
```

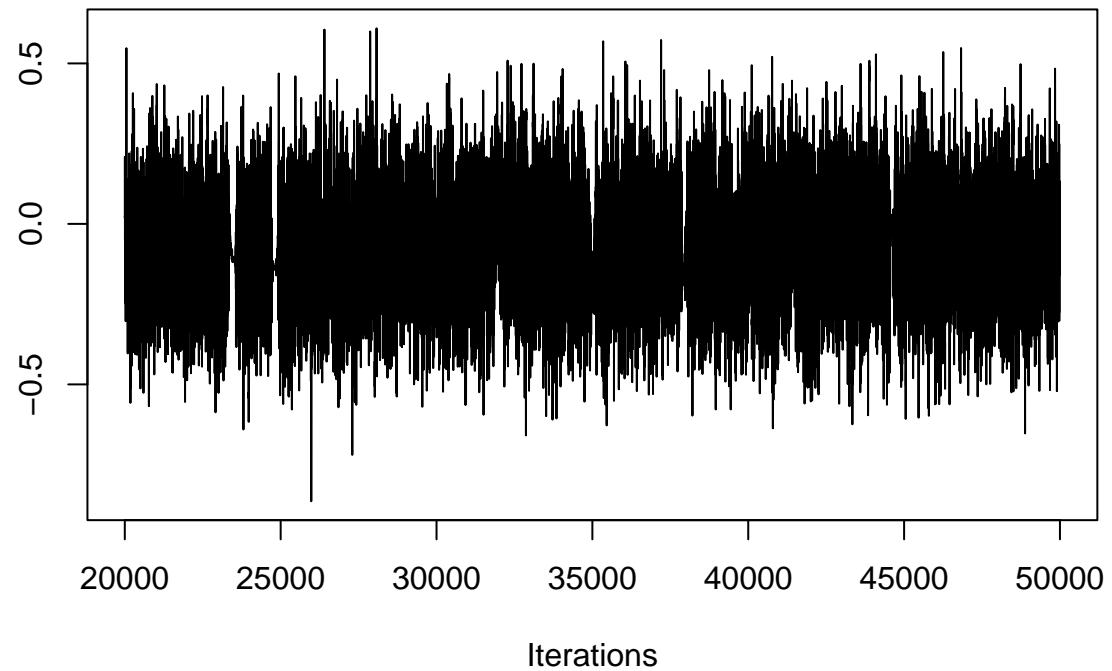
Trace of alpha[1]



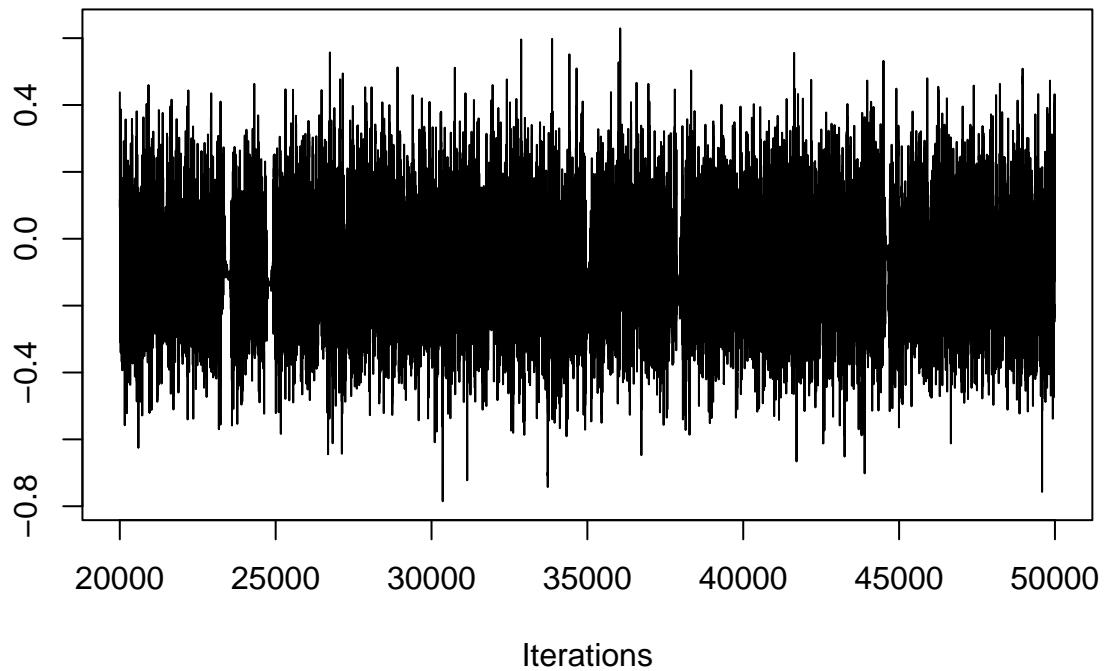
Trace of alpha[2]



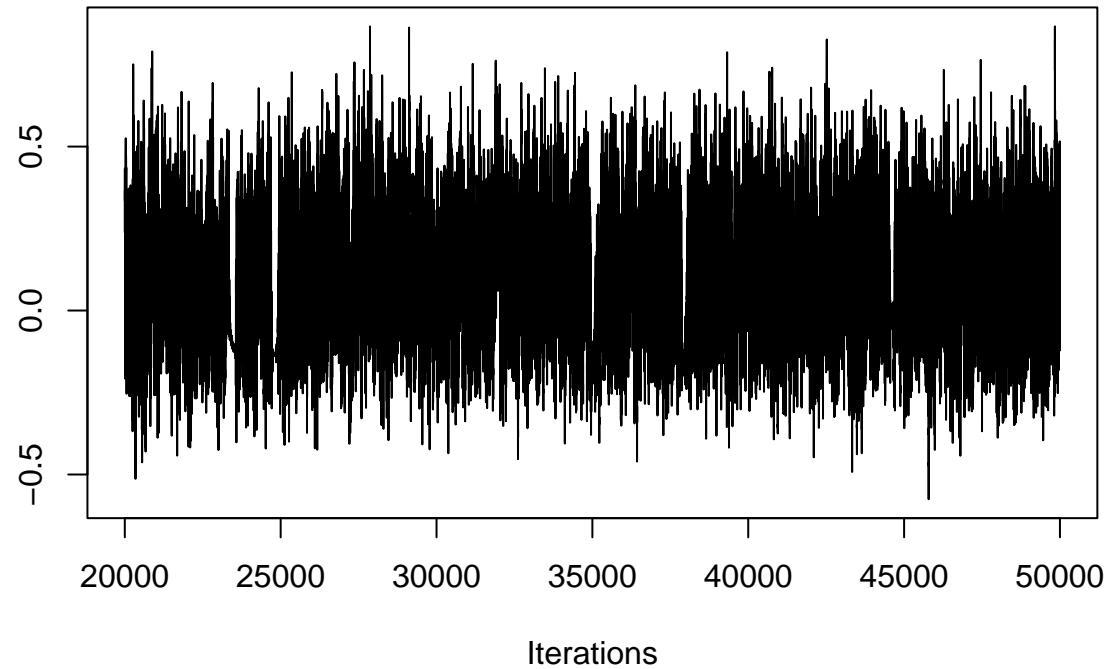
Trace of alpha[3]



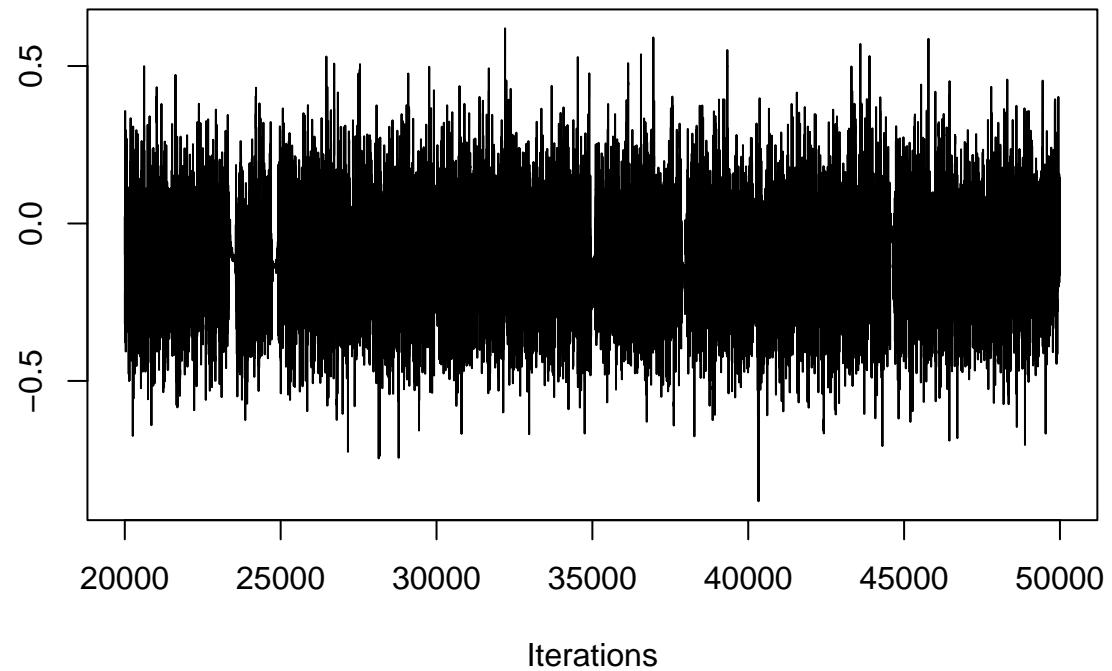
Trace of alpha[4]



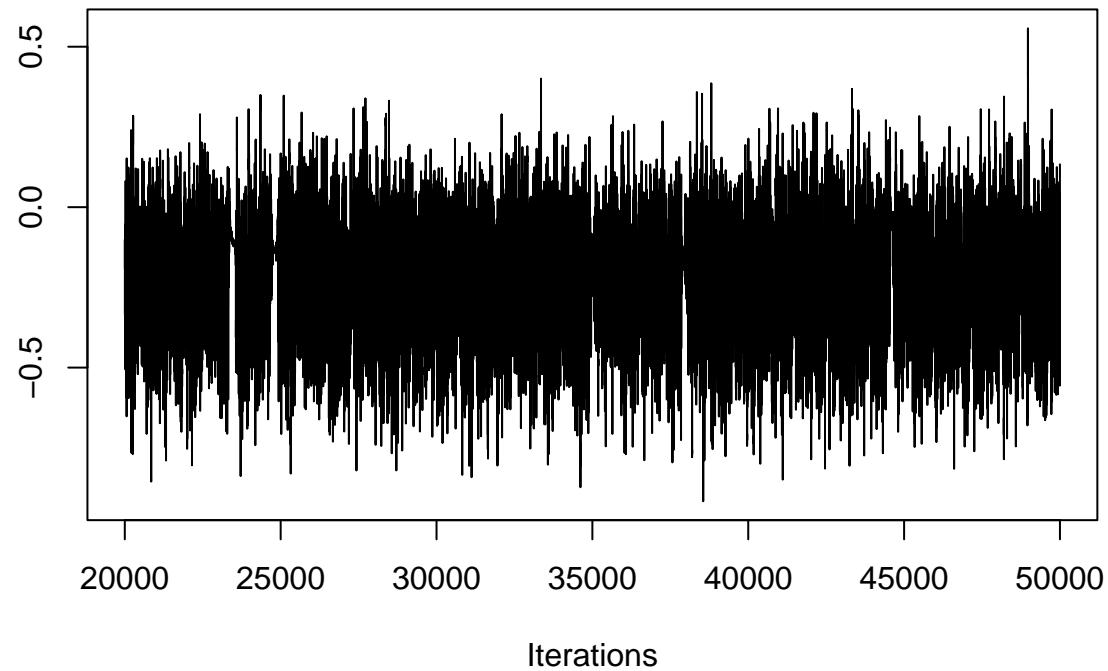
Trace of alpha[5]



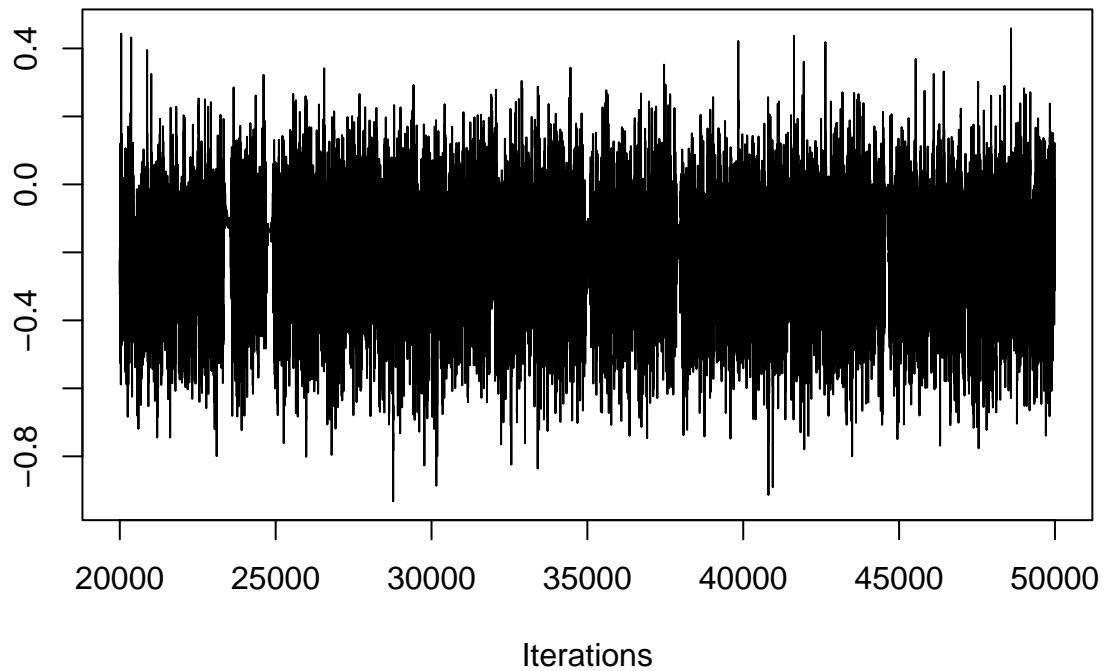
Trace of alpha[6]



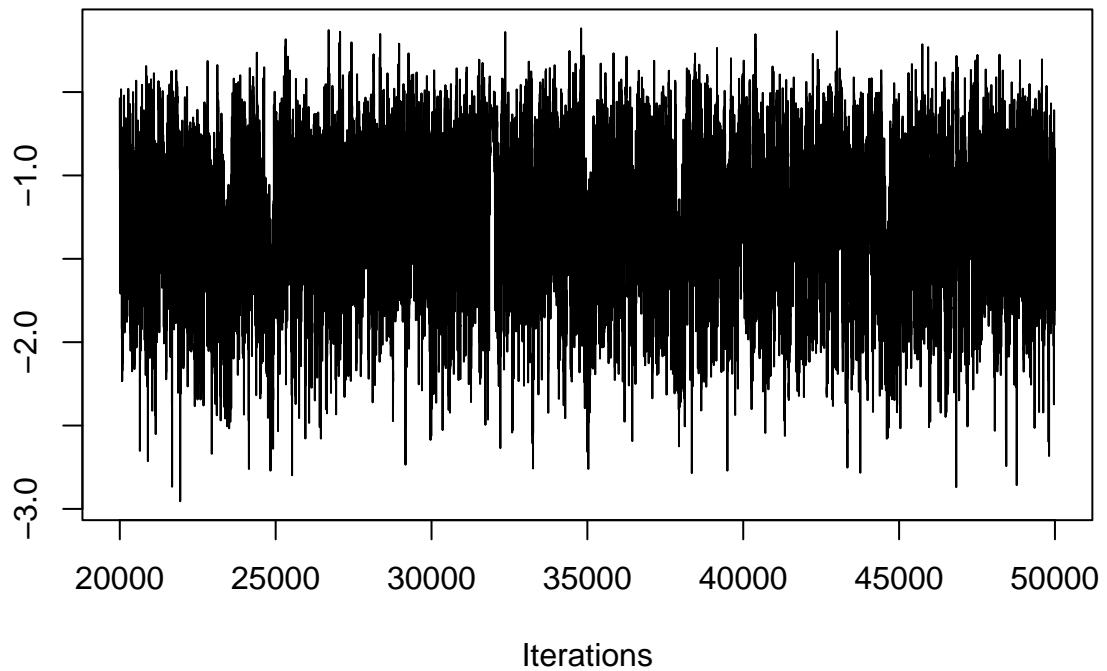
Trace of alpha[7]



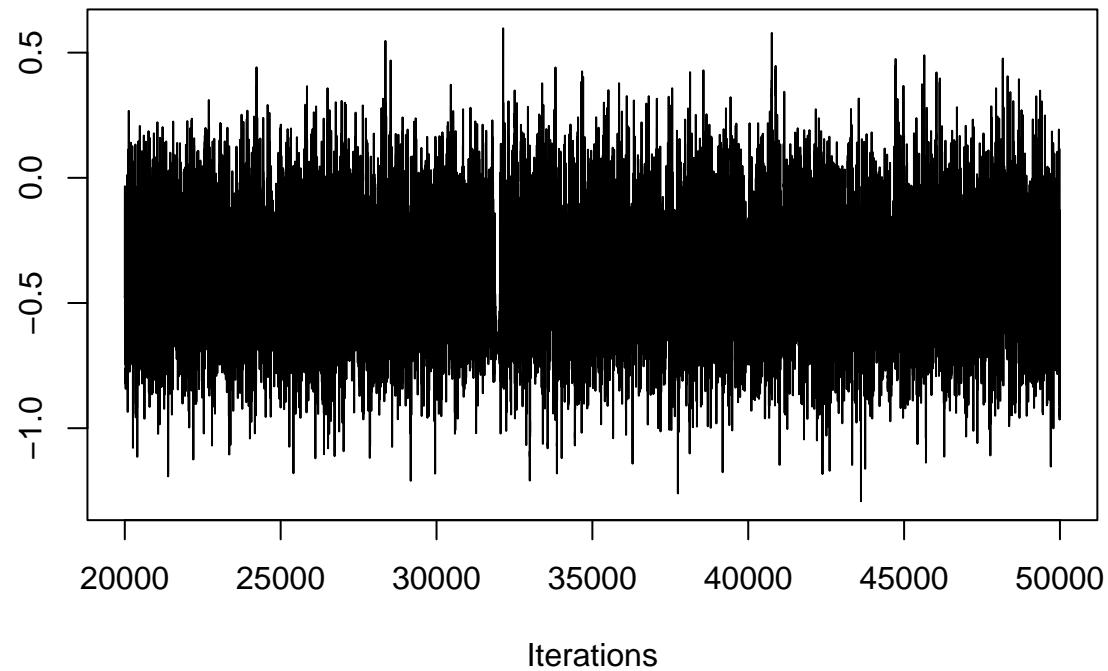
Trace of alpha[8]



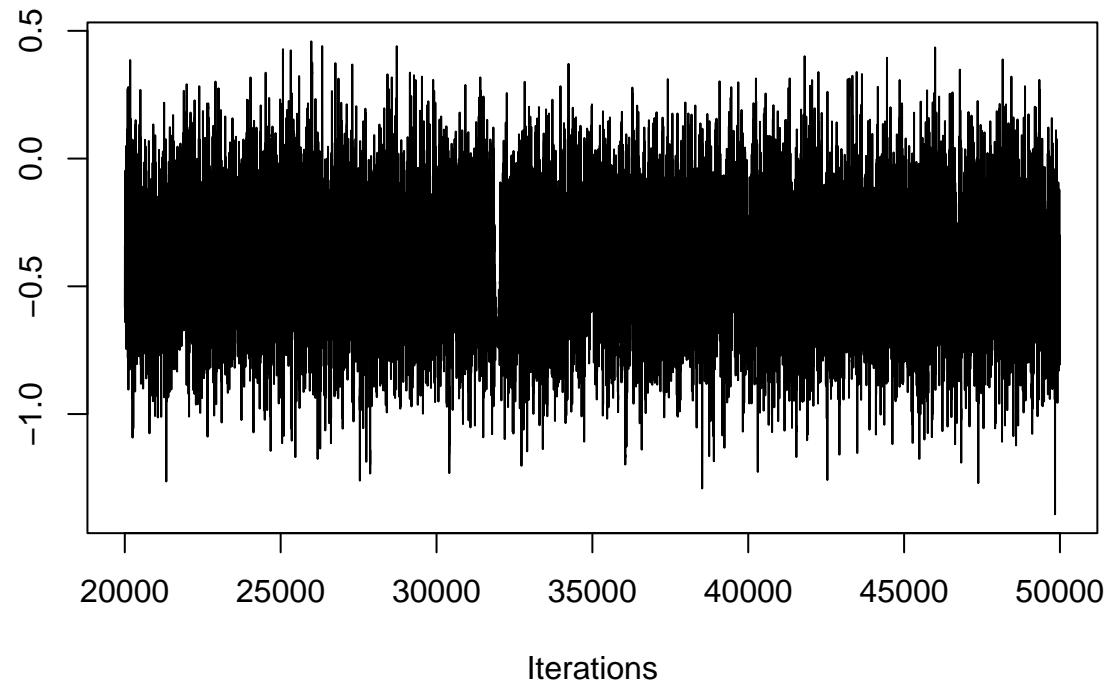
Trace of beta[1]



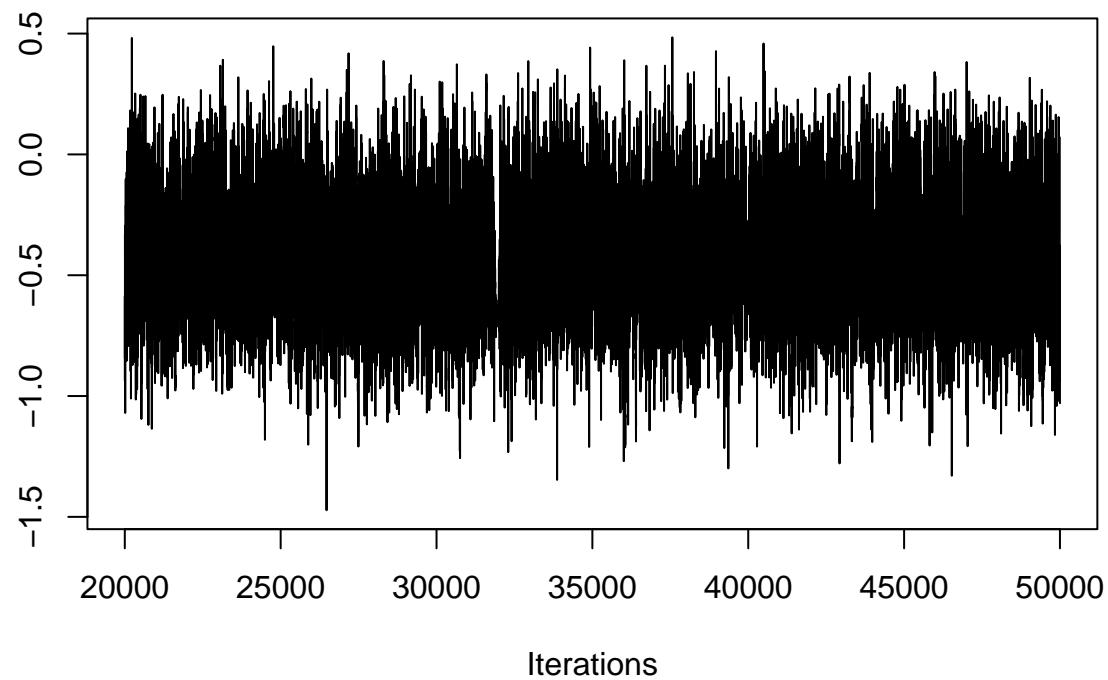
Trace of beta[2]



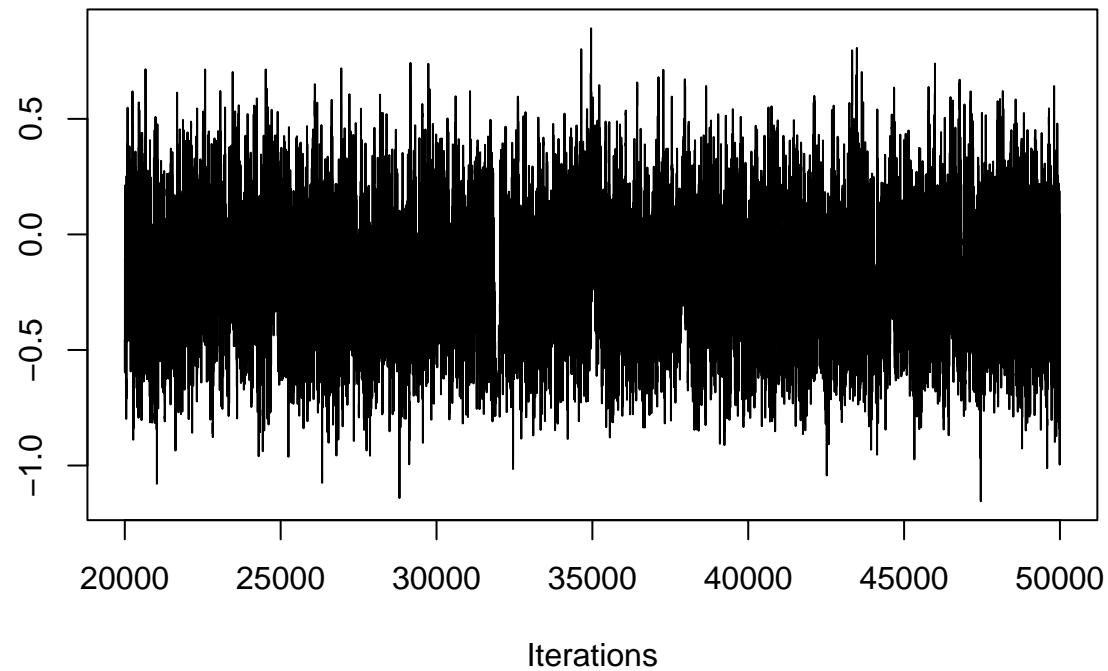
Trace of beta[3]



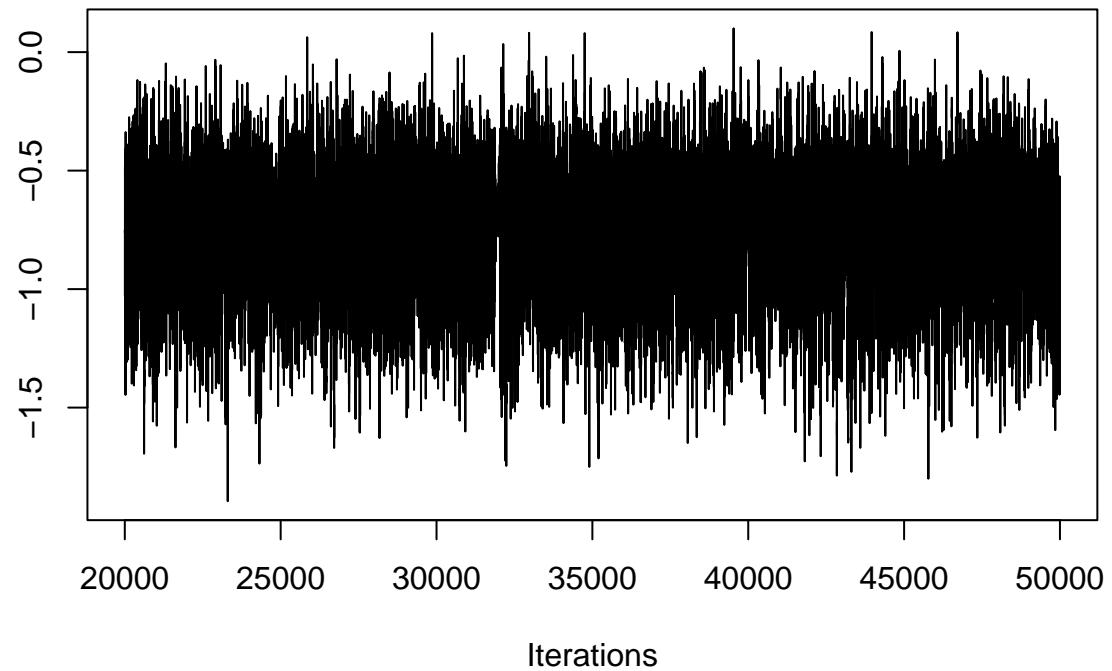
Trace of beta[4]



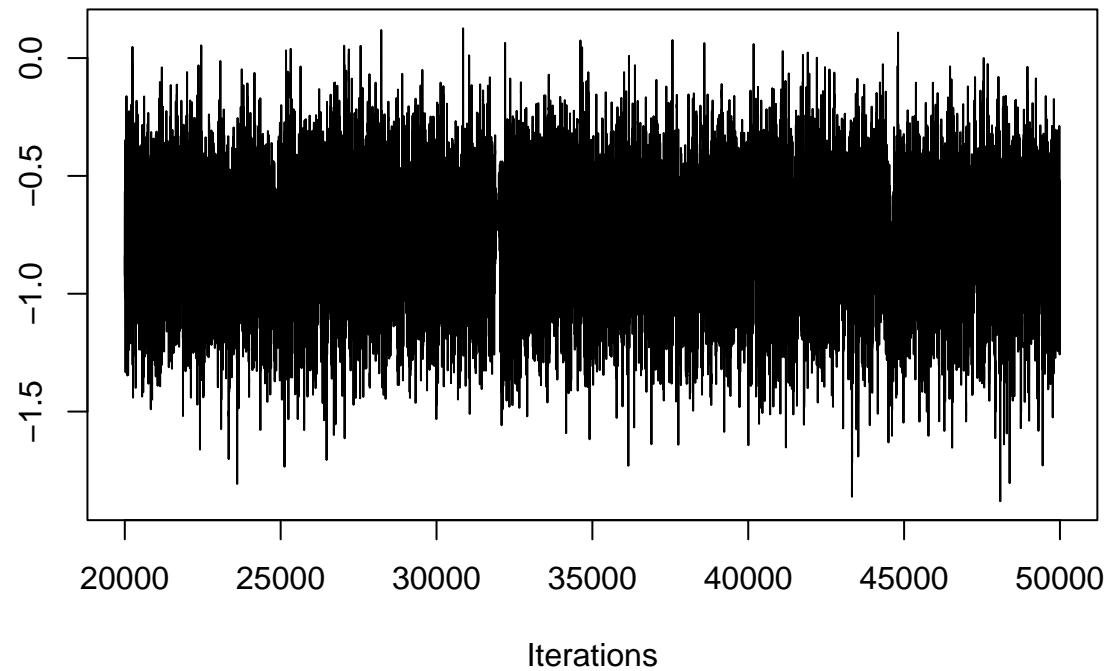
Trace of beta[5]



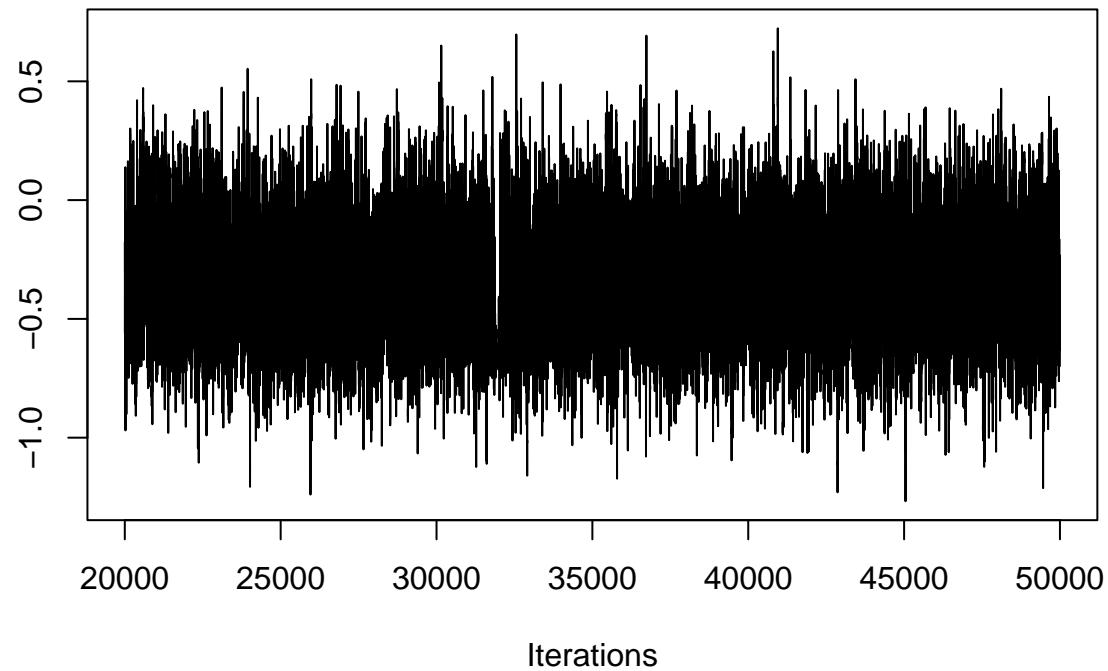
Trace of beta[6]



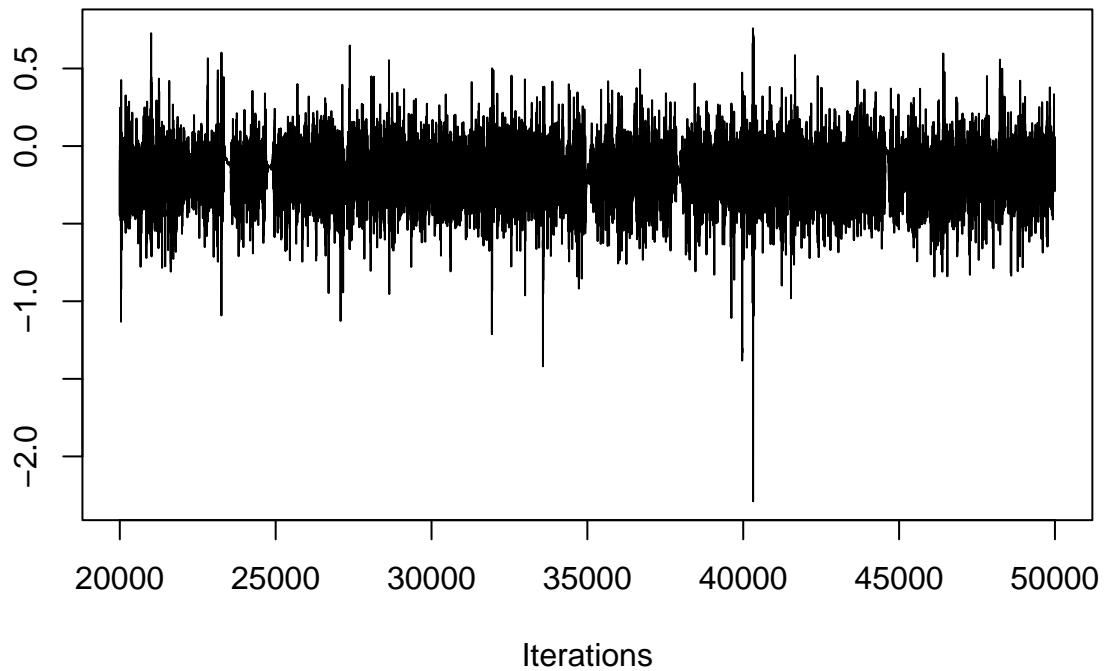
Trace of beta[7]



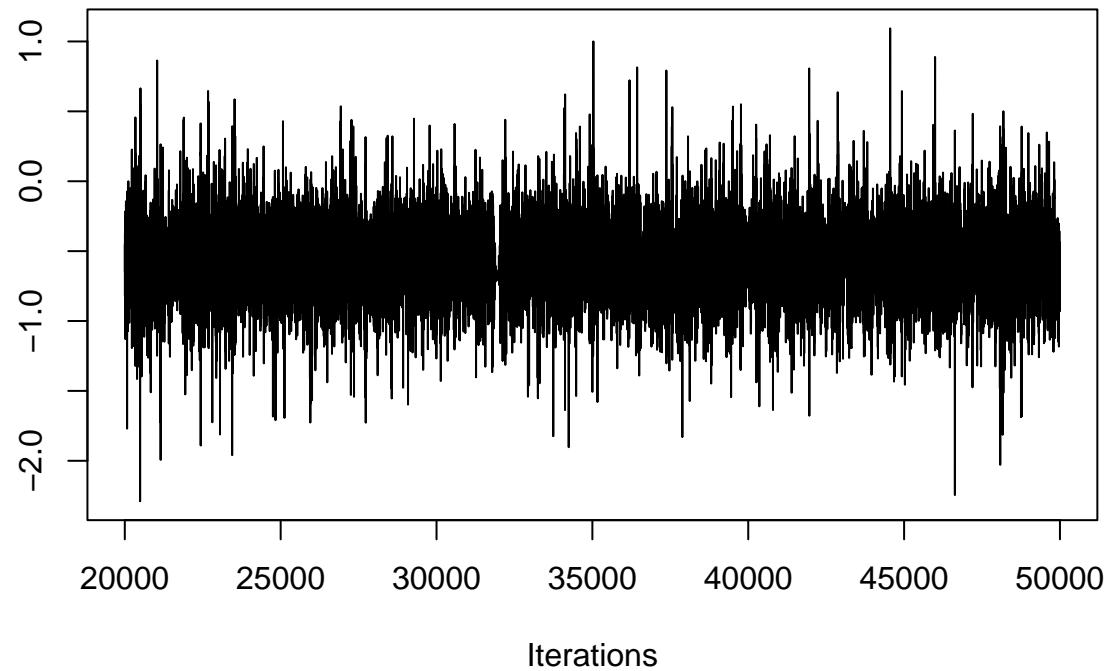
Trace of beta[8]



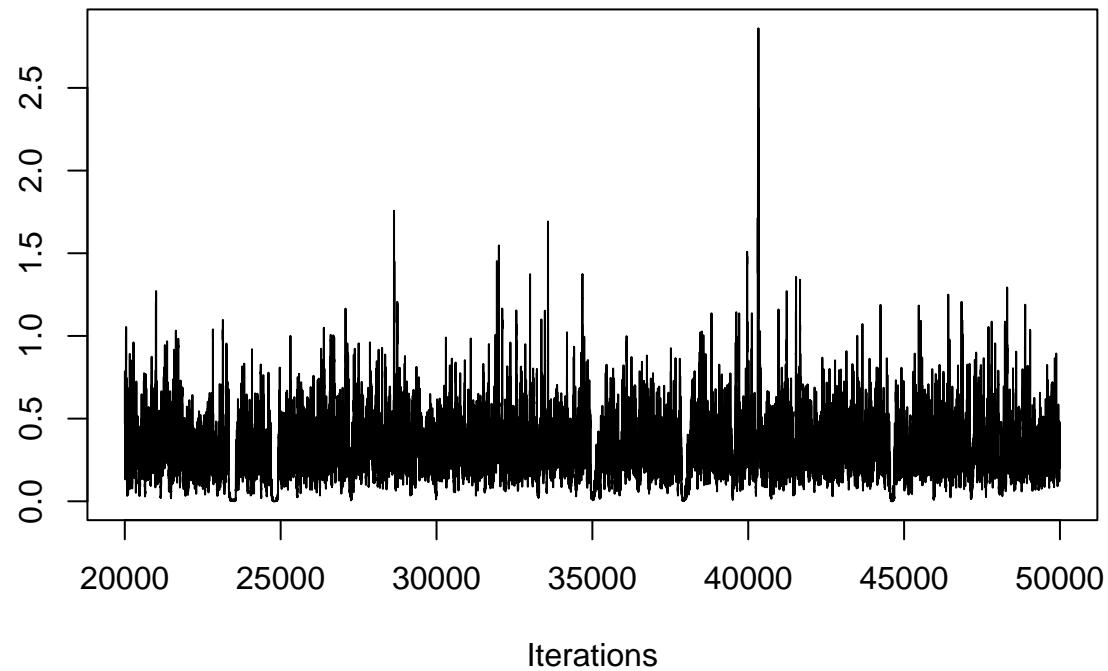
Trace of mu_alpha



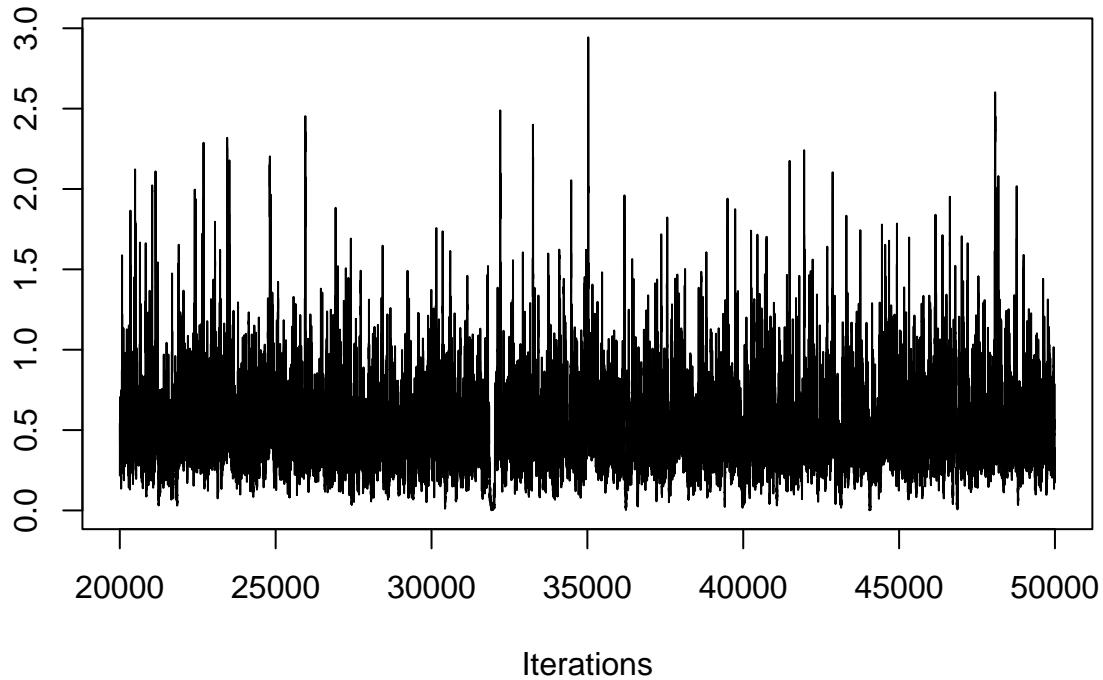
Trace of mu_beta



Trace of σ_{α}



Trace of sigma_beta



b

```
library(ggplot2)
library(tidyr)
library(dplyr)

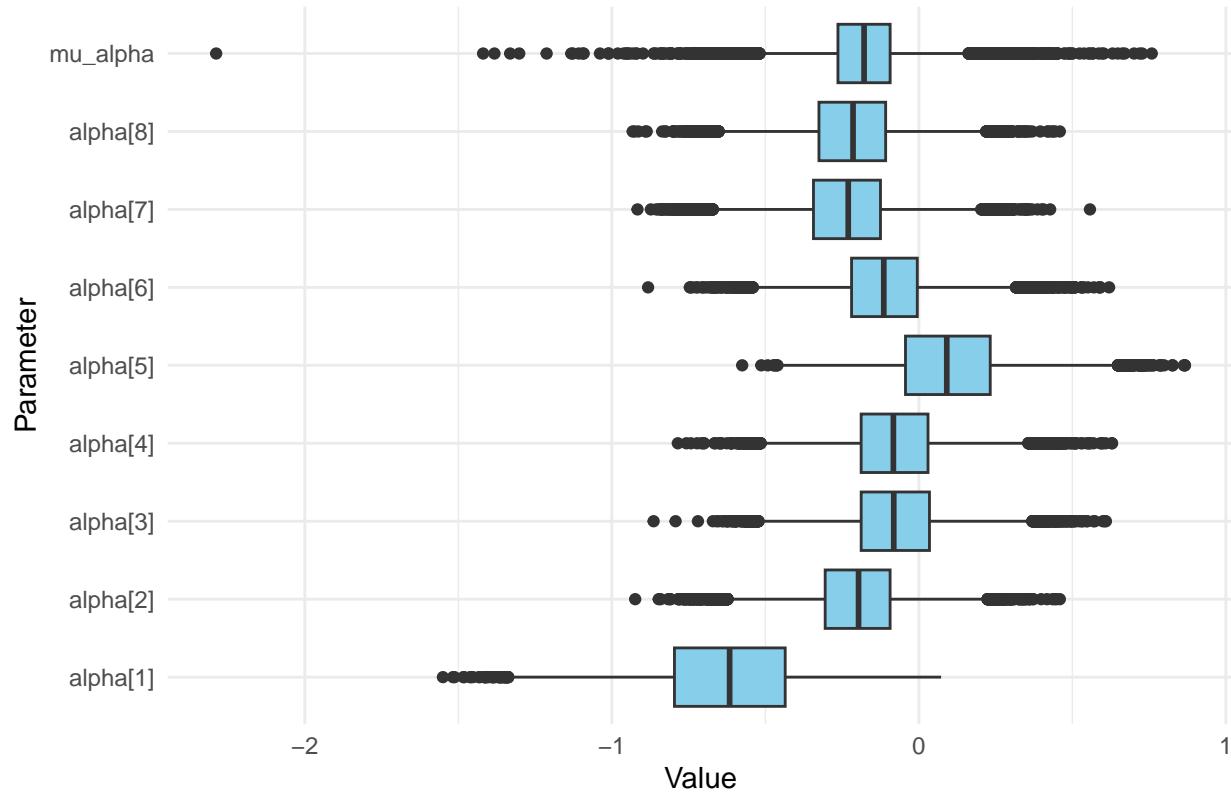
# Convert MCMC samples to data frame
samples_df <- as.data.frame(as.matrix(fit.samples))

# ----- Boxplot for alpha[j] and mu_alpha -----
alpha_cols <- grep("^\alpha\\\[", names(samples_df), value = TRUE)

alpha_long <- samples_df %>%
  select(all_of(alpha_cols), mu_alpha) %>%
  pivot_longer(cols = everything(), names_to = "Parameter", values_to = "Value")

ggplot(alpha_long, aes(x = Parameter, y = Value)) +
  geom_boxplot(fill = "skyblue") +
  theme_minimal() +
  coord_flip() +
  labs(title = "Posterior Boxplots: alpha[j] and mu_alpha", y = "Value", x = "Parameter")
```

Posterior Boxplots: alpha[j] and mu_alpha

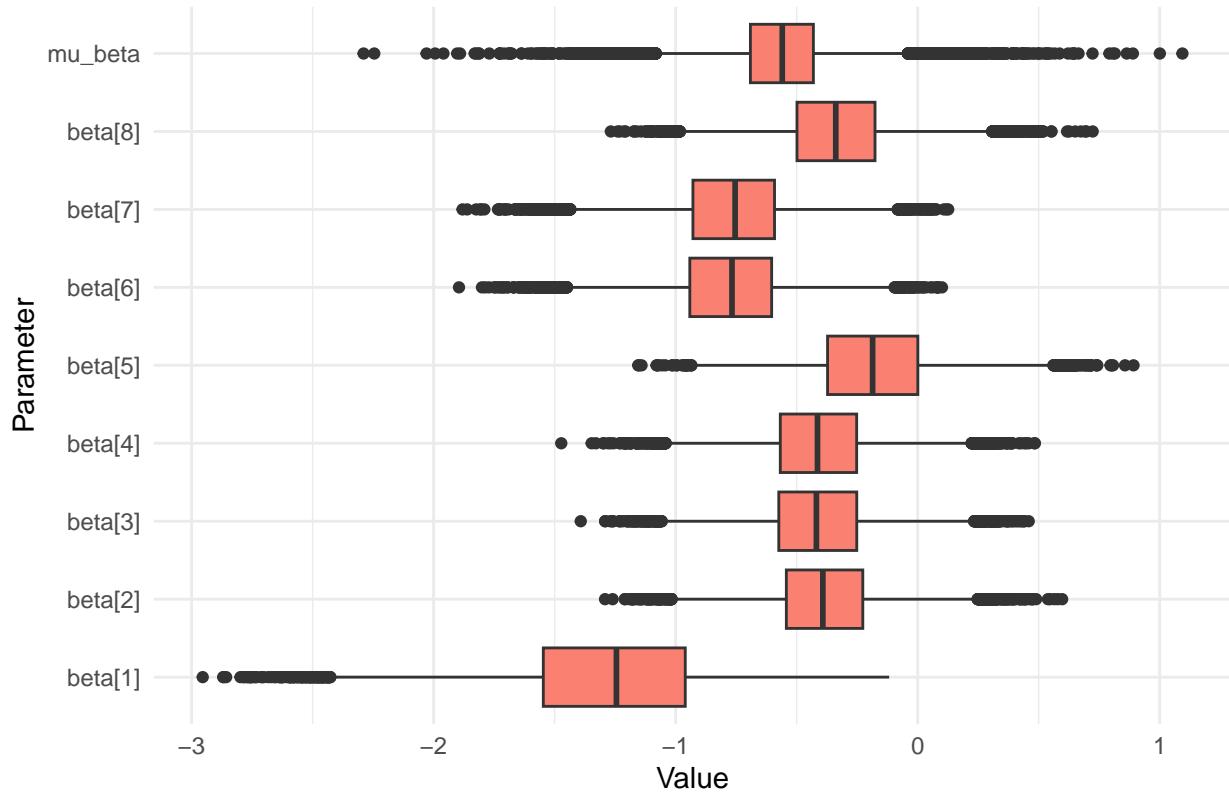


```
# ----- Boxplot for beta[j] and mu_beta -----
beta_cols <- grep("beta\\[\\]", names(samples_df), value = TRUE)

beta_long <- samples_df %>%
  select(all_of(beta_cols), mu_beta) %>%
  pivot_longer(cols = everything(), names_to = "Parameter", values_to = "Value")

ggplot(beta_long, aes(x = Parameter, y = Value)) +
  geom_boxplot(fill = "salmon") +
  theme_minimal() +
  coord_flip() +
  labs(title = "Posterior Boxplots: beta[j] and mu_beta", y = "Value", x = "Parameter")
```

Posterior Boxplots: beta[j] and mu_beta



C

```

set.seed(123)

# Inverse logit function
invlogit <- function(x) {
  1 / (1 + exp(-x))
}

# Convert samples to matrix
samples_df <- as.data.frame(as.matrix(fit.samples))

# Get posterior draws of mu_alpha, mu_beta, sigma_alpha, sigma_beta
N_samples <- nrow(samples_df)
theta_same <- numeric(N_samples)
theta_diff <- numeric(N_samples)

for (i in 1:N_samples) {
  # Simulate alpha* and beta* for a new school
  alpha_star <- rnorm(1, samples_df$mu_alpha[i], samples_df$sigma_alpha[i])
  beta_star <- rnorm(1, samples_df$mu_beta[i], samples_df$sigma_beta[i])

  # Predict friendship probability
}

```

```

    theta_same[i] <- invlogit(alpha_star + beta_star * 1)
    theta_diff[i] <- invlogit(alpha_star + beta_star * 0)
}

# Final results
mean(theta_same) # Predicted probability for x = 1

## [1] 0.338367

mean(theta_diff) # Predicted probability for x = 0

## [1] 0.456396

```

Problme 2

a

1. α_i : Baseline mean change in ΔPCV for cats of breed i , after adjusting for dose and type.
2. β_1 : Overall shift in ΔPCV across all cats, when breed, dose, type, and domestic status are held constant.
3. β_2 : Effect of dose on ΔPCV (expected change per unit increase in dose).
4. β_3 : Effect of disease type (Type = 1 for primary, 0 for secondary) on ΔPCV .
5. μ_{ij} : Expected ΔPCV for the j th cat of breed i .
6. σ^2 : Residual variance in ΔPCV across individual cats.
7. γ : Difference in breed-level baseline ΔPCV between domestic ($= 1$) and non-domestic ($= 0$) cats.
8. τ_α^2 : Variance in breed-specific intercepts α_i across all breeds.

b

```

cats <- read.table('~/Users/landlee/Desktop/STAT 3303/HW10/cats.txt',
                    header = TRUE)

cats <- cats %>%
  mutate(
    breed_index = as.numeric(factor(Breed, levels = c("A", "B", "C", "D", "E"))),
    domestic = ifelse(Breed %in% c("A", "B", "C"), 1, 0)
  )

# MCMC parameters
niter <- 10000
nburn <- 1000
nadapt <- 1000

# Number of breeds and total observations

```

```

Nb <- length(unique(cats$Breed))
N <- nrow(cats)

# Prepare JAGS data
data_jags <- list(
  y = cats$DeltaPCV,
  type = cats$Type,
  dose = cats$Dose,
  breed = cats$breed_index,
  domestic = c(1,1,1,0,0),
  Nb = Nb,
  N = N
)

# Initial values
inits <- list(
  beta1 = 0,
  beta2 = 0,
  beta3 = 0,
  gamma = 0,
  tau = 1,
  tau_alpha = 1,
  alpha = rep(0, Nb)
)

model <- "model {
  for (i in 1:N) {
    y[i] ~ dnorm(mu[i], tau)
    mu[i] <- alpha[breed[i]] + beta1 + beta2 * dose[i] + beta3 * type[i]
  }

  for (j in 1:Nb) {
    alpha[j] ~ dnorm(gamma * domestic[j], tau_alpha)
  }

  beta1 ~ dnorm(0, 1/100)
  beta2 ~ dnorm(0, 1/100)
  beta3 ~ dnorm(0, 1/100)
  gamma ~ dnorm(0, 1/100)
  tau ~ dgamma(3, 10)
  tau_alpha ~ dgamma(3, 10)

  sigma <- 1 / sqrt(tau)
}""

fit <- jags.model(textConnection(model),
  data = data_jags,
  inits = inits,
  n.chains = 1,
  n.adapt = nadapt)

## Compiling model graph
## Resolving undeclared variables

```

```

##      Allocating nodes
## Graph information:
##      Observed stochastic nodes: 40
##      Unobserved stochastic nodes: 11
##      Total graph size: 228
##
## Initializing model

params <- c("beta1", "beta2", "beta3", "gamma", "tau", "tau_alpha", "alpha", "sigma")
cat.samples <- coda.samples(fit, variable.names = params, n.iter = niter)
summary(cat.samples)

##
## Iterations = 1:10000
## Thinning interval = 1
## Number of chains = 1
## Sample size per chain = 10000
##
## 1. Empirical mean and standard deviation for each variable,
##    plus standard error of the mean:
##
##           Mean        SD   Naive SE Time-series SE
## alpha[1] -2.8542 1.37009 0.0137009     0.125250
## alpha[2] -1.3676 1.35136 0.0135136     0.129822
## alpha[3] -2.8312 1.33598 0.0133598     0.140356
## alpha[4]  0.3676 1.27734 0.0127734     0.123278
## alpha[5] -0.5469 1.28079 0.0128079     0.123330
## beta1    -1.3596 1.38097 0.0138097     0.149997
## beta2    -0.3899 0.09021 0.0009021     0.003651
## beta3    0.2428 0.37526 0.0037526     0.007299
## gamma   -2.3372 1.65047 0.0165047     0.129536
## sigma    1.1327 0.13101 0.0013101     0.001657
## tau      0.8101 0.18209 0.0018209     0.002227
## tau_alpha 0.4054 0.19229 0.0019229     0.003683
##
## 2. Quantiles for each variable:
##
##           2.5%       25%       50%       75%     97.5%
## alpha[1] -5.7474 -3.706496 -2.8293 -1.9179 -0.3065
## alpha[2] -4.2479 -2.214109 -1.3626 -0.4341  1.1510
## alpha[3] -5.6908 -3.669838 -2.8101 -1.9253 -0.3786
## alpha[4] -2.3734 -0.402899  0.3957  1.2240  2.6867
## alpha[5] -3.2498 -1.339166 -0.5164  0.3378  1.7429
## beta1    -3.8454 -2.310527 -1.4022 -0.4998  1.5853
## beta2    -0.5647 -0.450291 -0.3902 -0.3289 -0.2105
## beta3    -0.4979 -0.007459  0.2416  0.4952  0.9780
## gamma   -5.7262 -3.338920 -2.3103 -1.2619  0.8444
## sigma    0.9118  1.040459  1.1198  1.2118  1.4242
## tau      0.4930  0.680980  0.7975  0.9237  1.2029
## tau_alpha 0.1233  0.264557  0.3742  0.5162  0.8530

# Convert MCMC samples to matrix
mcmc_matrix <- as.matrix(cat.samples)

```

```

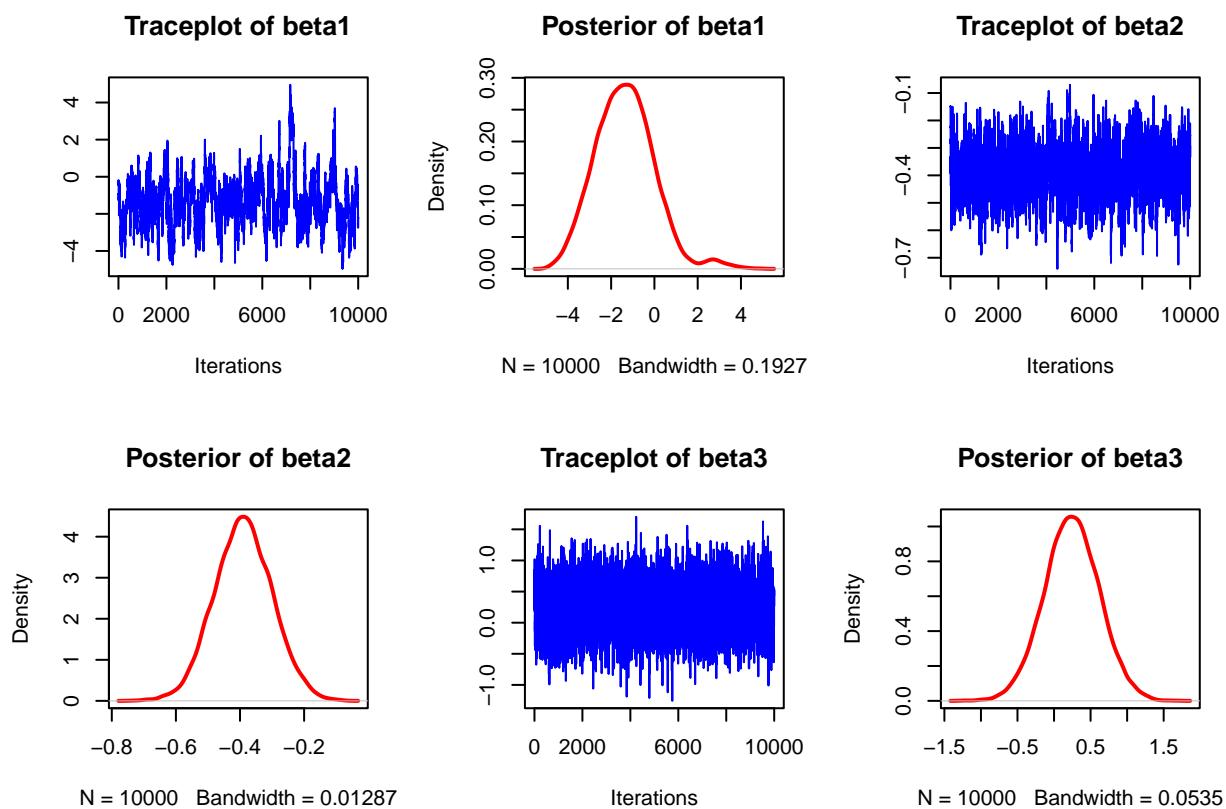
# Parameters to plot
params <- c("beta1", "beta2", "beta3", "gamma", "tau", "tau_alpha")

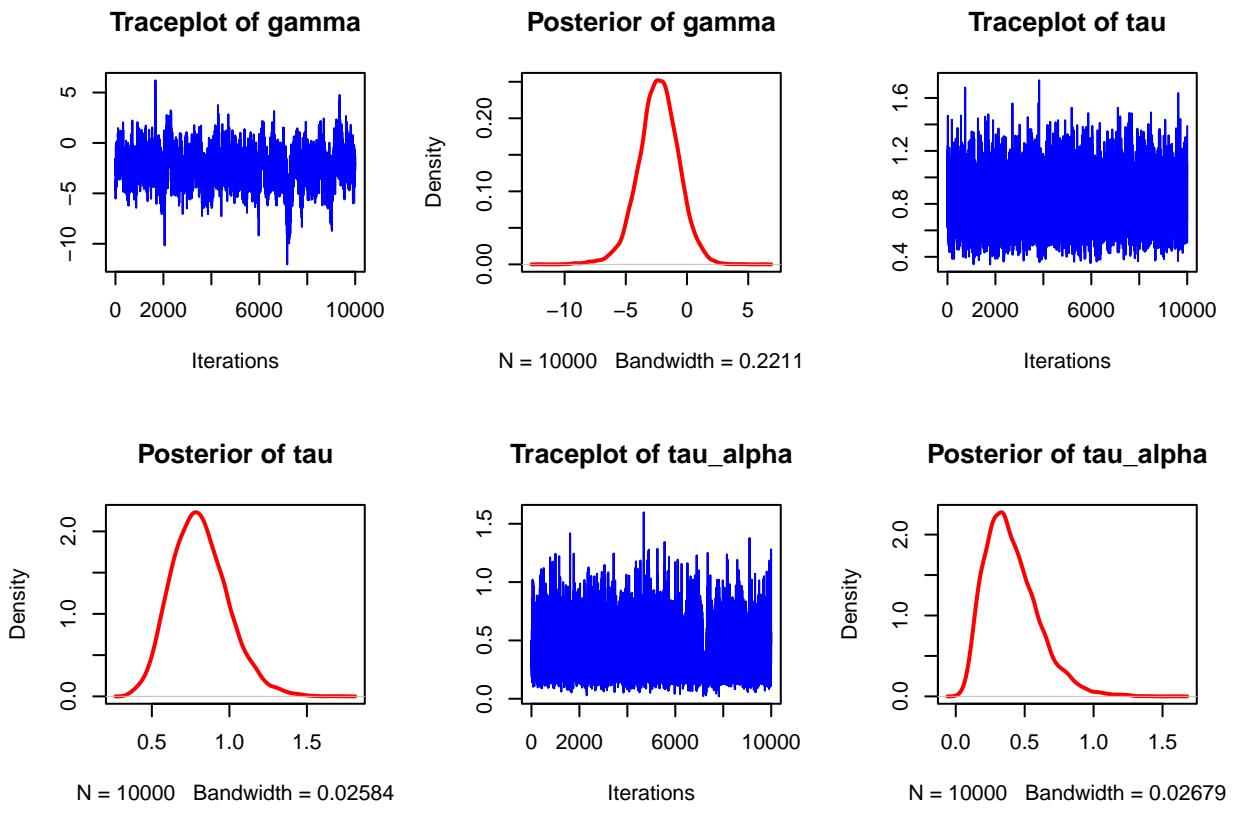
# Set layout: 2 rows, 3 columns
par(mfrow = c(2, 3))

# Loop over each parameter
for (param in params) {
  # Traceplot for one parameter at a time
  traceplot(cat.samples[, param],
            main = paste("Traceplot of", param),
            col = "blue", lwd = 1)

  # Density plot
  plot(density(mcmc_matrix[, param]),
        main = paste("Posterior of", param),
        col = "red", lwd = 2)
}

```



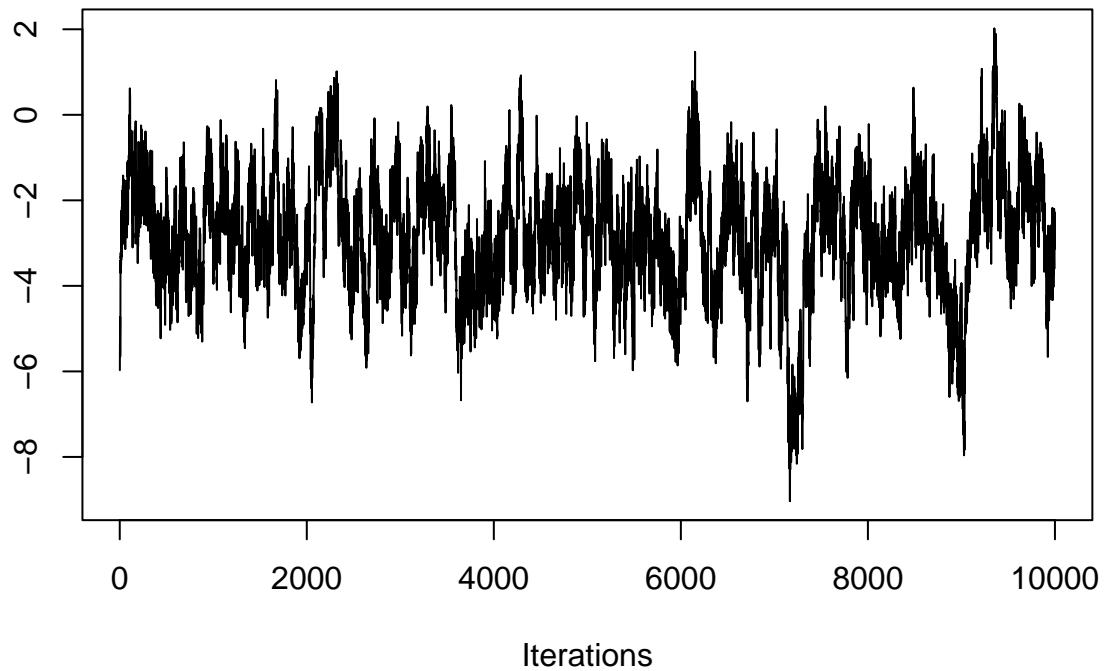


```
ess <- effectiveSize(cat.samples)
print(round(ess, 2))
```

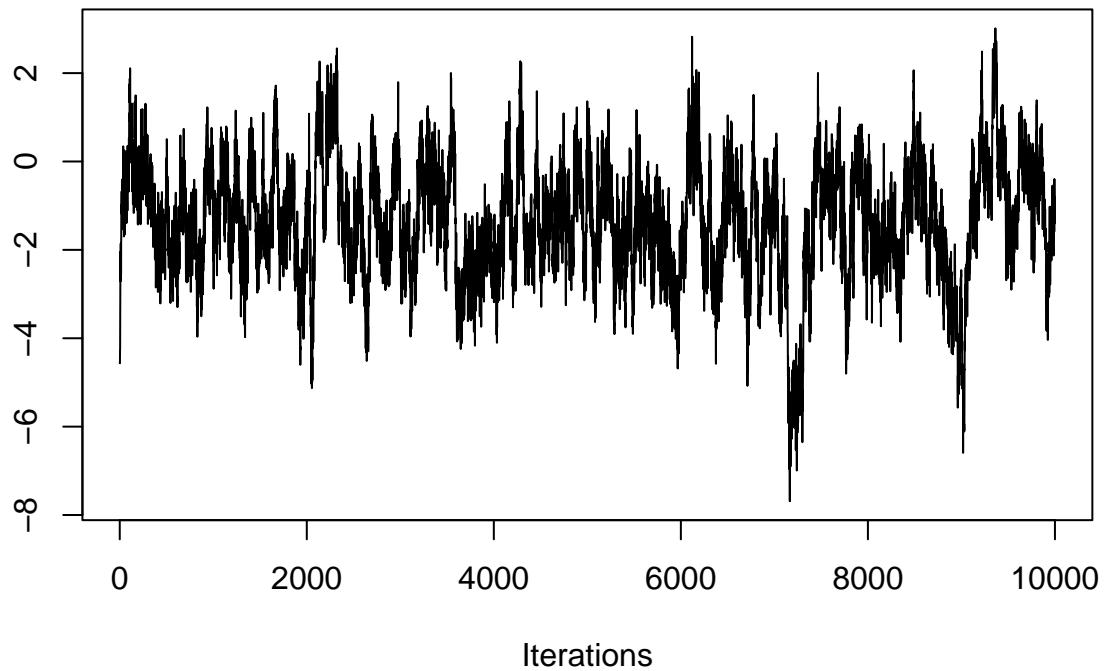
```
##   alpha[1]   alpha[2]   alpha[3]   alpha[4]   alpha[5]     beta1     beta2     beta3
##   119.66    108.35    90.60    107.36   107.85    84.76   610.65  2643.02
##   gamma      sigma     tau  tau_alpha
##   162.34    6249.89  6683.84   2726.16
```

```
traceplot(cat.samples)
```

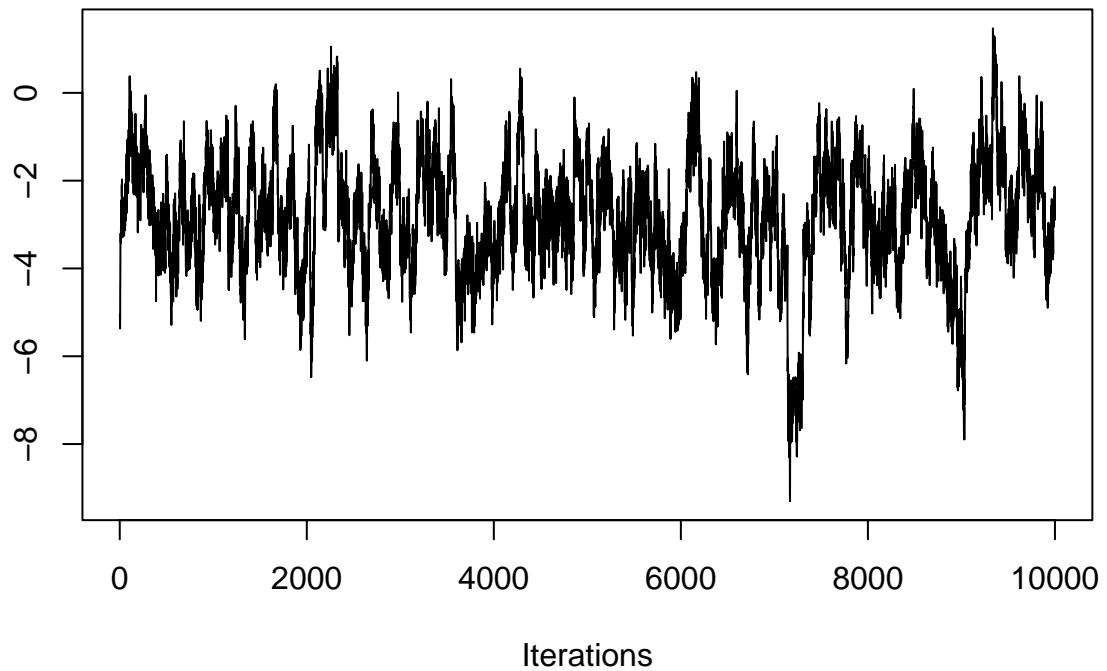
Trace of alpha[1]



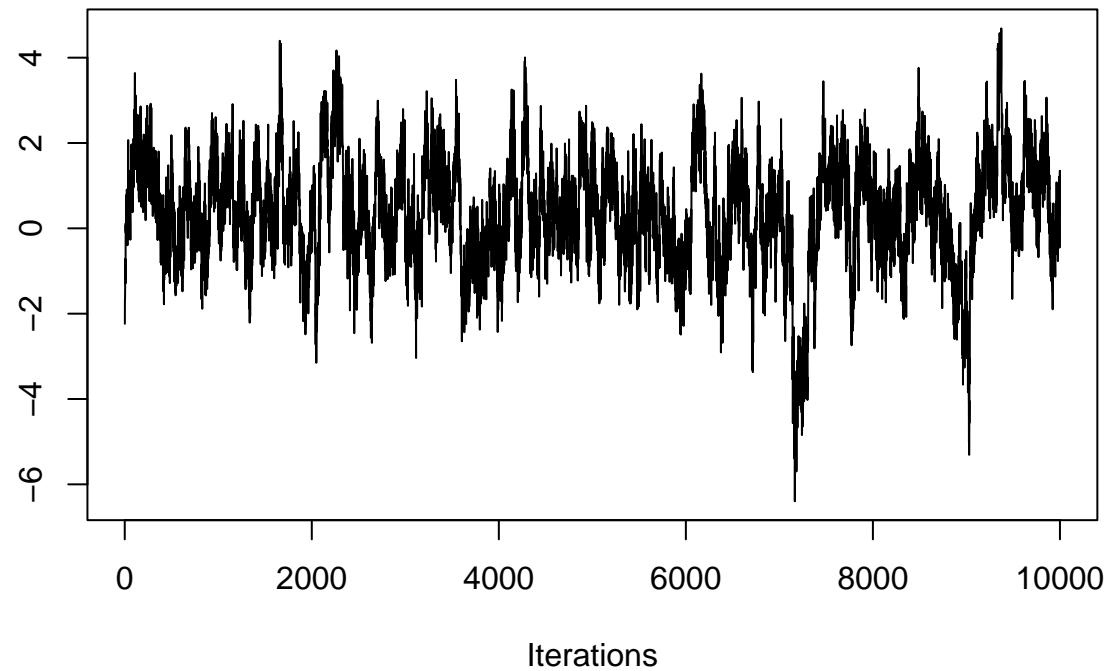
Trace of alpha[2]



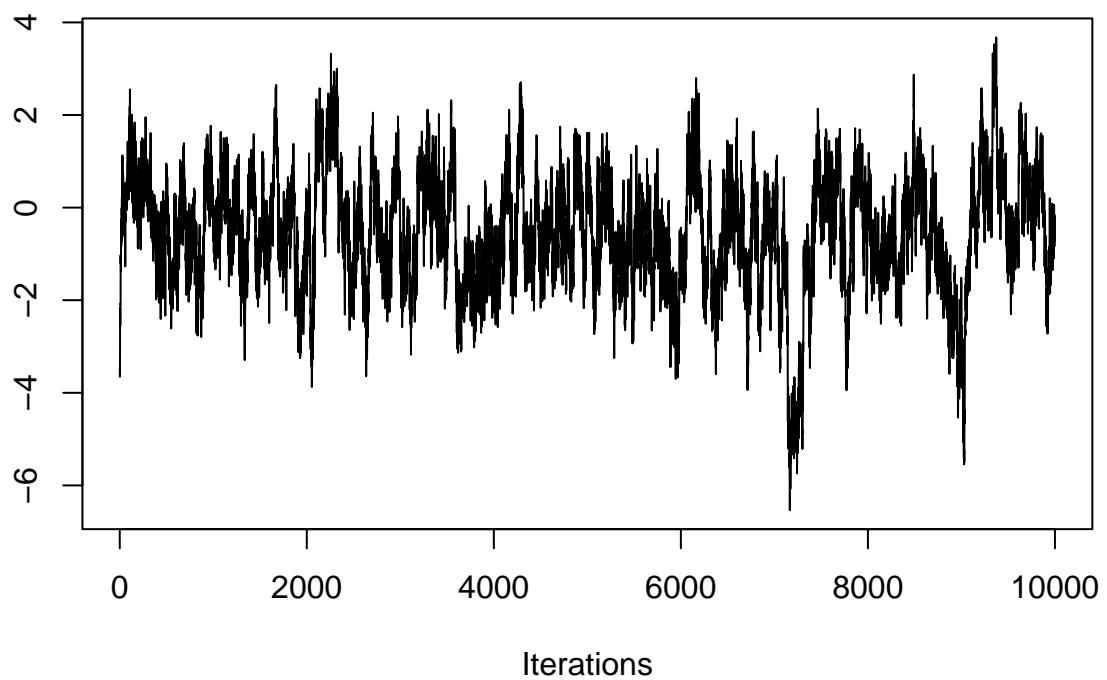
Trace of alpha[3]



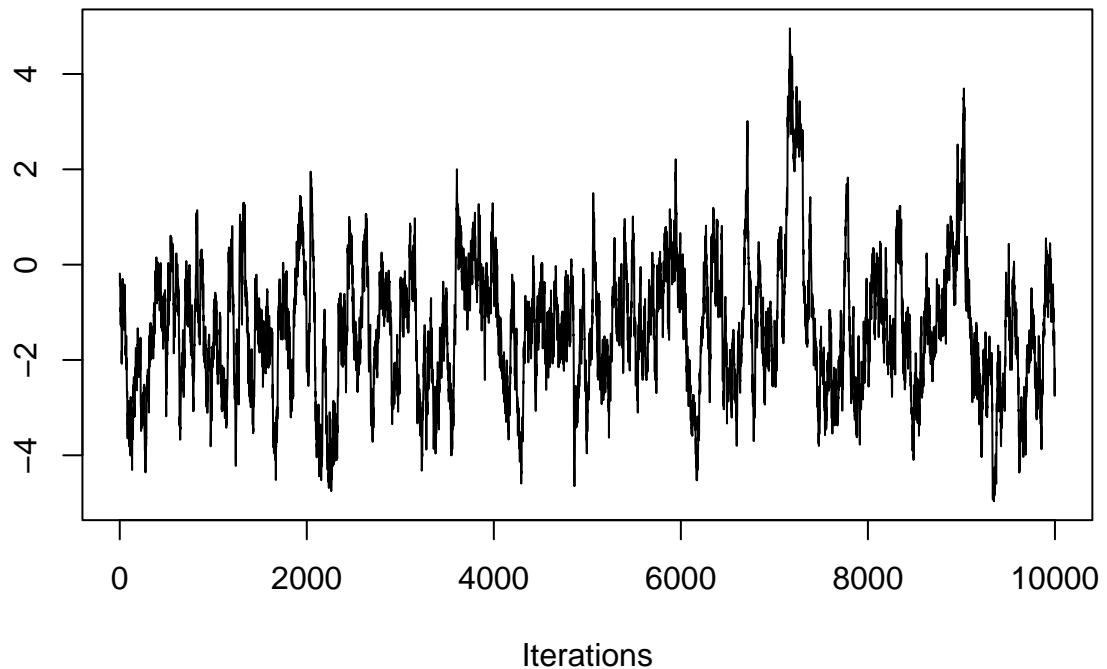
Trace of alpha[4]



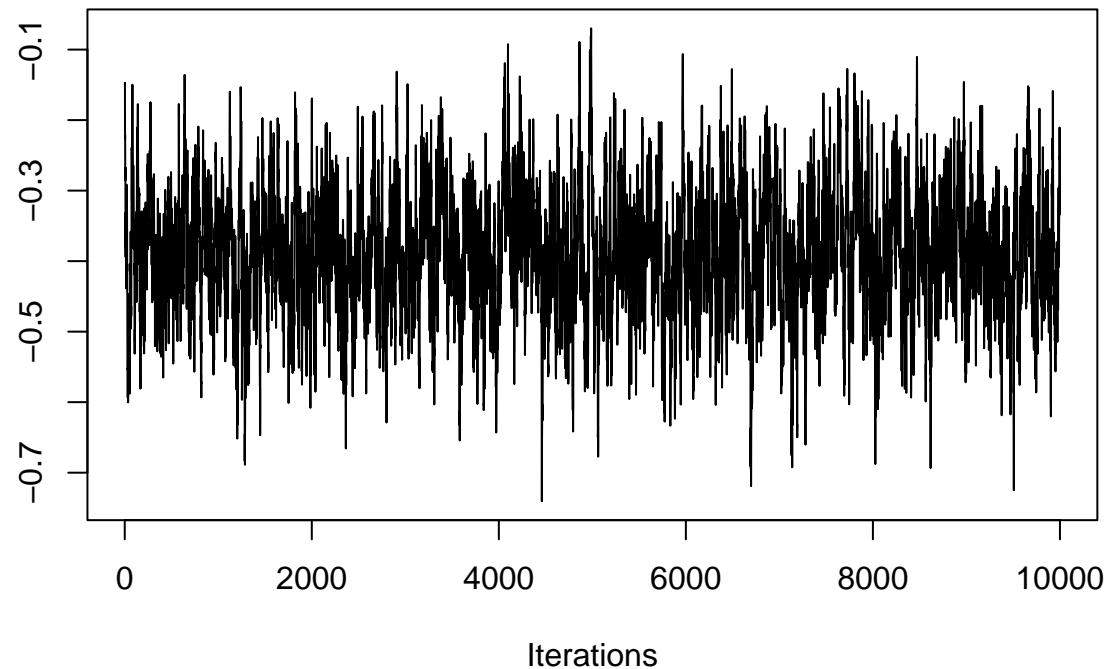
Trace of alpha[5]



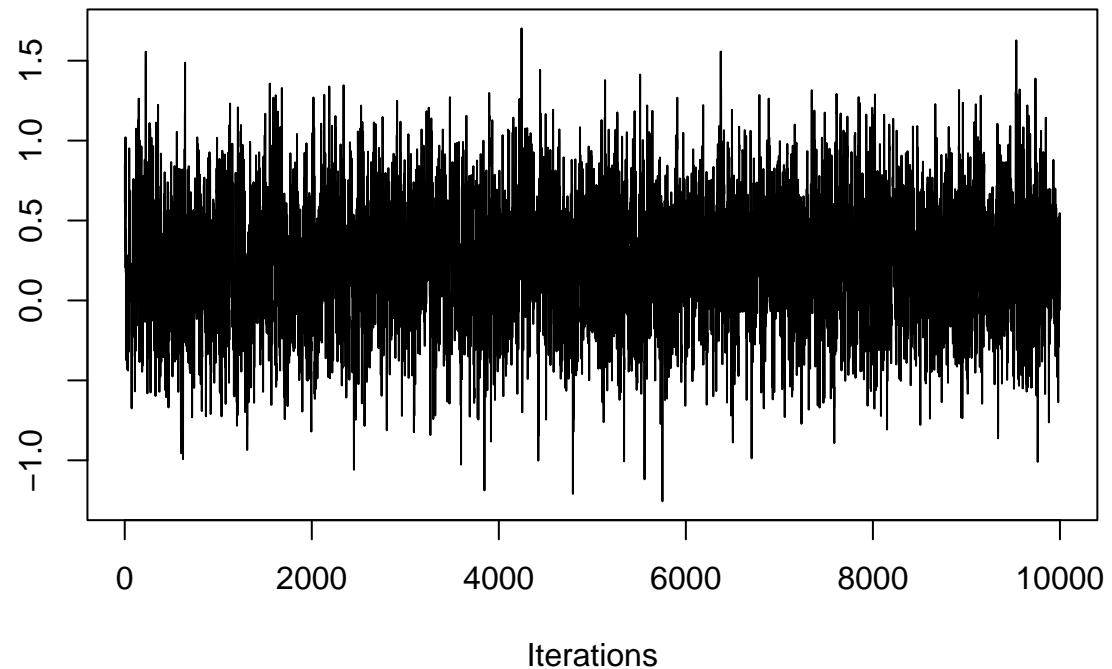
Trace of beta1



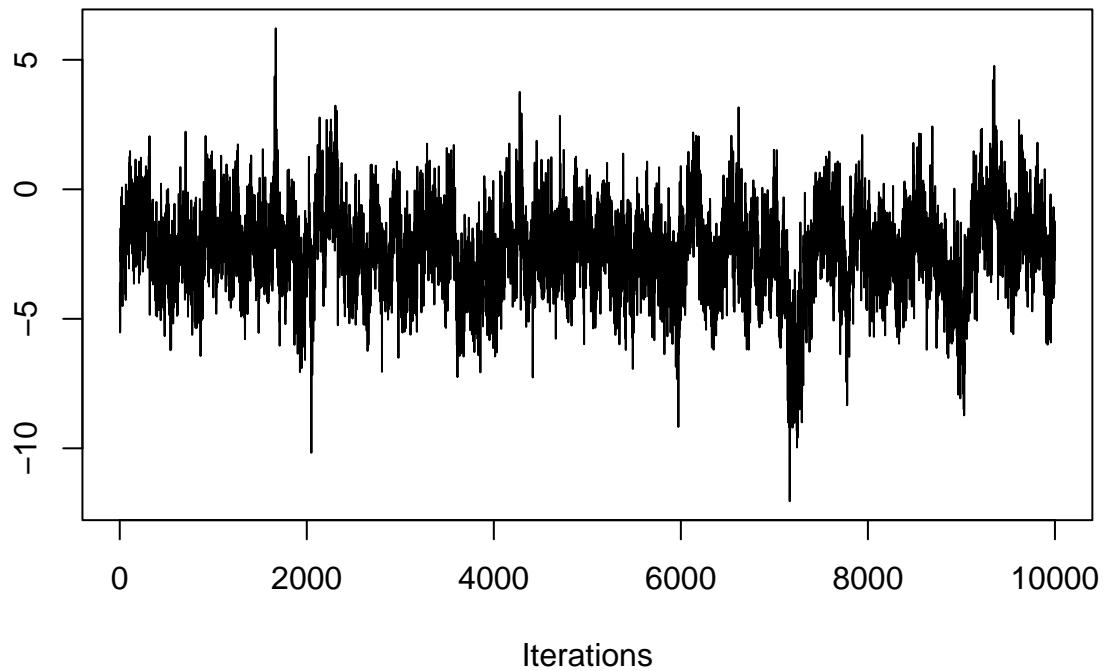
Trace of beta2



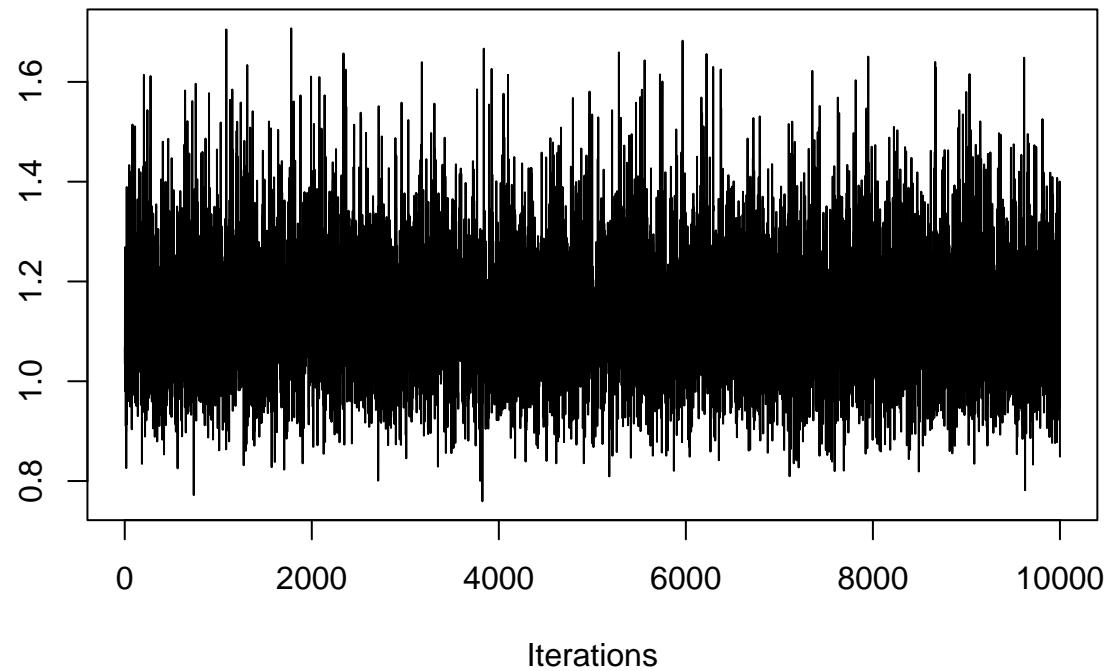
Trace of beta3



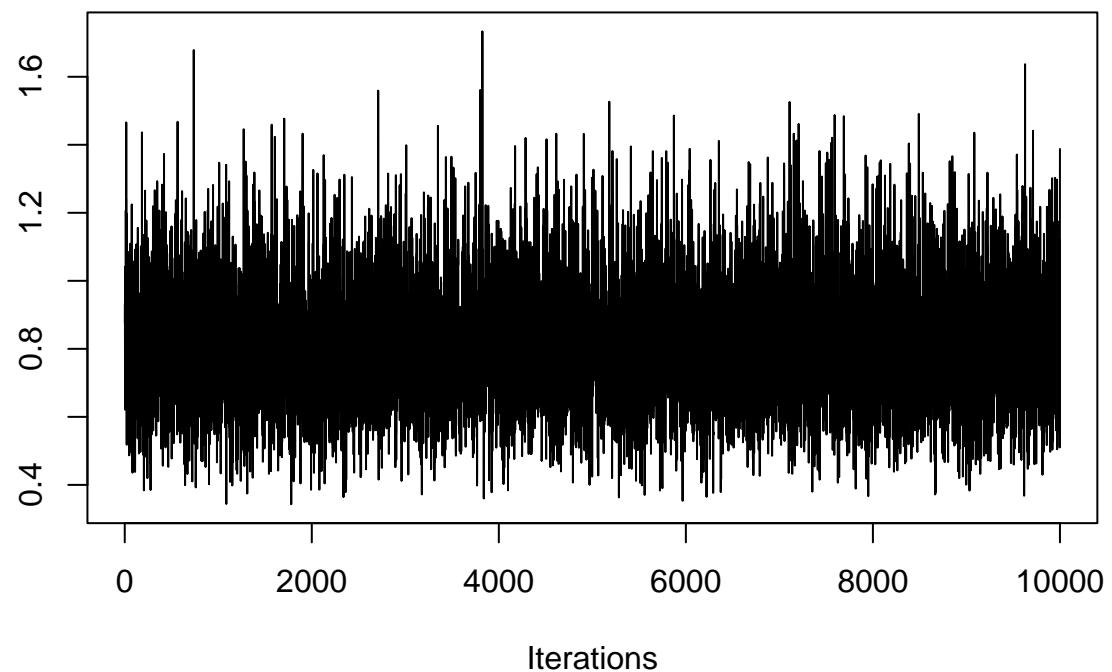
Trace of gamma



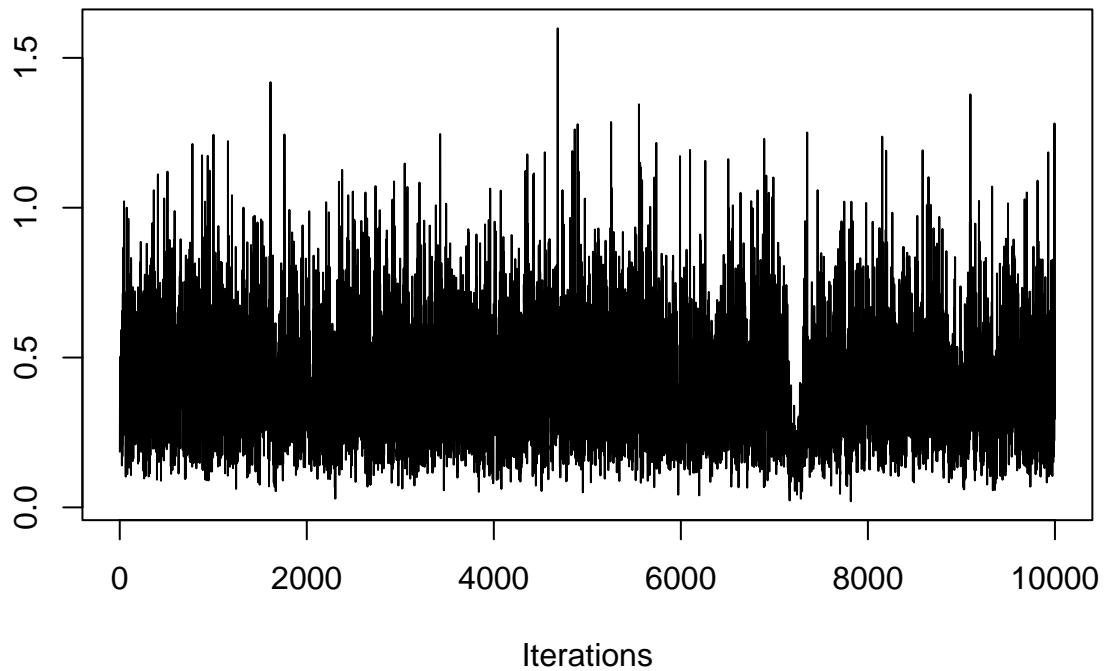
Trace of sigma



Trace of tau



Trace of tau_alpha



c

Mean	SD	Naive SE	Time-series SE
------	----	----------	----------------

beta2	-0.3928	0.08698	0.0008698	0.003443	2.5%	25%	50%	75%	97.5%	beta2	-0.5591	-0.45095	-0.3946	-0.3365
	-0.2162													

Yes. $\beta_2 = -0.3960$ and the confidence interval doesn't include 0

d

Mean	SD	Naive SE	Time-series SE
------	----	----------	----------------

gamma	-2.3542	1.56515	0.0090364	0.0630482	2.5%	25%	50%	75%	97.5%	gamma	-5.6342	-3.32079	-2.3110
	-1.3267	0.7726											

No. Because the confidence interval contains 0.

e

The measure of uncertainty I used in both part (c) and part (d) is the 95% posterior credible interval, based on the 2.5% and 97.5% quantiles of the posterior distributions which can be found in the summary.